

8 Tree-Based Models

Andrew Liang

11/22/2020

Notes

- Useful for interpretation, but typically not competitive with best supervised learning approaches in terms of prediction

Basics of Decision Trees

Regression Trees

Process of Building a Regression Tree:

1. Divide the predictor space (set of possible values of X_1, X_2, \dots, X_p) into J distinct and non-overlapping regions R_1, R_2, \dots, R_J
 2. For every observation in the R_j region, we make the same prediction, which is simply the mean of the response values for the training observations in R_j
- In step 1, we construct R_1, R_2, \dots, R_J such that the predictor space is divided into high-dimensional rectangles, or *boxes*, for ease of interpretation
 - goal is to minimize RSS given by:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

where \hat{y}_{R_j} is the mean response for the training observations within the j th box.

- Another way to look at it is that we consider all predictors X_1, X_2, \dots, X_p and all possible values of cutpoint s for each of the predictors, then choose the predictor and cutpoint such that the resulting tree has lowest RSS
 - for any j and s , we define pair of half-planes:

$$R_1(j, s) = \{X | X_j < s\} \quad \text{and} \quad R_2(j, s) = \{X | X_j \geq s\}$$

and we seek to value of j and s to minimize:

$$\sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

where \hat{y}_{R_1} is the mean response for training observations in $R_1(j, s)$ and \hat{y}_{R_2} is mean response for train obs in $R_2(j, s)$

Once the regions R_1, R_2, \dots, R_J have been created, we can then predict response for a given test observation using mean of train obs in the region to which that test obs belongs

Tree Pruning

- Process above may likely overfit data as the resulting tree may be too complex
 - smaller tree with fewer splits can lead to lower variance and better interpretation at the cost of a little bias
- Better strategy may be to grow a very large tree T_0 , and then *prune* it back to get a subtree
 - want to get a subtree that leads us to the lowest test error rate
 - however, using CV for every subtree may be infeasible, so we need to select a small set of subtrees
- can use *cost complexity pruning*, consider a sequence of trees indexed by a nonnegative tuning parameter α
 - for each value of α , there corresponds a subtree $T \subset T_0$ such that it minimizes:

$$\sum_{m=1}^{|T|} \sum_{i: x \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T|$$

where $|T|$ indicates the number of terminal nodes of tree T , and R_m is the region corresponding to the m th terminal node

- as the number of terminal nodes increases, there is a penalty α , so the above quantity will tend to be minimized for a smaller subtree
 - select α using CV

Algorithm for building Regression Trees

1. Use recursive binary splitting to grow a large tree on training data, stopping when each terminal node has fewer than some minimum # of observations
2. Apply cost complexity pruning to large tree in order to obtain a sequence of best subtrees as a function of α
3. Use K-fold CV to choose α . Divide training observations into K fold. For each $k = 1, \dots, K$:
 - repeat steps 1 and 2 on all but k th fold of training data
 - evaluate mean squared prediction error on data in left-out k th fold, as a function of α
 - average out the results for each value of α , and pick α to minimize error
4. Return the subtree from step 2 that corresponds to chosen value of α

Classification Trees

- very similar to regression tree, but predicts qualitative response instead
 - predict that each observation belongs to the *most commonly occurring class* or training observations in the region to which it belongs
- also interested in the *class proportions* among training observations that fall into that region
- also use recursive binary splitting to grow a classification tree, but instead of RSS, *classification error rate* is used as the criterion for making the binary splits
 - simply the fraction of training obs in that region that don't belong to the most common class:

$$E = 1 - \max_k(\hat{p}_{mk})$$

where \hat{p}_{mk} is the proportion of training observations in the m th region that are from the k th class. In practice however, two other measures are preferable:

- The *Gini index*:

$$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$

is a measure of total variance (of which we want to minimize) across the K classes. G takes on a small value if \hat{p}_{mk} is close to zero or one. It is a measure of *node purity* - a small value indicates that a node contains predominantly observations from a single class

- The *Entropy*:

$$D = - \sum_{k=1}^K \hat{p}_{mk} \log(\hat{p}_{mk})$$

similarly to the Gini, will take on a value near zero if all \hat{p}_{mk} are near zero or one. Both measurements are quite similar numerically

- generally, classification error rate is preferable if prediction accuracy is the goal of the final pruned tree

Trees vs Linear Models

- if relationship between features and response is well approximated by a linear model, then a linear regression would outperform trees
- if there is a highly non-linear and complex relationship between features and response, then trees may outperform linear regression

Pros of Trees

- very easy to explain
- may closely mirror human decision-making more than regression and classification approaches in previous methods
- displayed graphically and easily interpreted
- can handle qualitative predictors without the need to create dummy variables

Cons of Trees

- generally don't have same level of predictive accuracy as other regression and classification techniques
- tend to be very non-robust, small change in data can cause large change in the tree

Bagging, Random Forests, Boosting

Bagging

- in order to solve the high variance problem in trees, *bagging* can help reduce it through bootstrapping procedures
 - recall in bootstrap, given n independent observations Z_1, \dots, Z_n , each with variance σ^2
 - * variance of mean \bar{Z} is given by $\frac{\sigma^2}{n}$, thus has a lower variance
 - * essentially we use repeated samples from our original training data to create B different training sets:

$$\hat{f}_{bag}(x) = \frac{1}{B} \sum_{b=1}^B \hat{f}^{*b}(x)$$

which is called bagging

- in the context of regression trees, we construct B different trees using B bootstrapped training sets, then avg the resulting predictions
 - avg out all the trees reduces the variance
- in classification trees, for a given test observation, we can record the class predicted by each of the B trees, and take a *majority vote*, which is the overall prediction most commonly occurring among all the B predictions
- important to note that using a large B will not lead to overfitting
 - generally use $B = 100$ to achieve sufficient performance

Out-of-Bag Error Estimation

- very straightforward way of estimating test error of a bagged model, without the need of CV
- on avg, each bagged tree makes use around two-thirds of observations (sampling with replacement of training set)
 - remaining one-third of observations not used are the *Out-of-Bag* (OOB) observations
 - can predict response for i th observation using each of the trees in which that observation was OOB
 - * yields around $B/3$ predictions for i th observation
 - with B sufficiently large, OOB error is essentially the same as LOOCV
 - * convenient when perform CV would be computationally infeasible

Variable Importance Measures

- bagging improves prediction accuracy at the expense of interpretability
 - can obtain summary of important predictors using RSS or Gini index
 - * for regression, record the total amount the RSS is decreased due to splits over a given predictor
 - * for classification, sum the total amount that the Gini index is decreased by splits over a given predictor

Random Forests

- improves over bagging through decorrelation of the trees
- similar to bagging, we build a number of trees based on bootstrapping, but now we choose a *random sample of m predictors* as split candidates from the full set of p predictors
 - split is only allowed to use one of those m predictors
 - * a fresh sample of m predictors chosen at each split
 - typically choose $m \approx \sqrt{p}$
 - thus at each split, the algorithm is not even allowed to consider a majority of the available predictors
 - * helps *decorrelate* trees, to prevent the domination of one strong predictor on all the trees
 - when $m = p$, then the process is just bagging, hence bagging is a special case of random forest

Boosting

- boosting works similarly to bagging, but now each tree is grown *sequentially*
 - uses info from previous tree
 - doesn't use bootstrap, instead tree is fit on a modified version of original data

Algorithm for Boosting

1. Set $\hat{f}(x) = 0$ and $r_i = y_i$ for all i in the training set
2. For $b = 1, 2, \dots, B$, repeat:
 - Fit a tree \hat{f}^b with d splits ($d + 1$ terminal nodes) to training data (X, r)
 - Update \hat{f} by adding shrunk version of new tree:

$$\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$$

- Update the residuals:

$$r_i \leftarrow r_i + \lambda \hat{f}^b(x)$$

3. Output the boosted model:

$$\hat{f}(x) = \sum_{b=1}^B \lambda \hat{f}^b(x)$$

- the idea is to slowly improve \hat{f} in areas where it doesn't perform well
- shrinkage λ slows the process down, allowing more and different shaped trees to attack residuals

Boosting has three tuning parameters:

1. Number of trees B . Unlike bagging and random forests, boosting can overfit if B is too large, hence we use CV to choose B

2. shrinkage parameter λ , a small positive number and controls rate at which boosting learns (typical values are 0.01 or 0.001). Very small λ can require large B to achieve good performance
3. Number of splits d in each tree. Controls complexity of boosted ensemble, and often $d = 1$ works well, where each tree is a *stump* of a single split. When $d = 1$, boosted ensemble is fitting an additive model, since each term involves only one variable. Generally, d is the *interaction depth*
 - because growth of a tree depends on previous trees, smaller trees are typically sufficient
 - smaller trees can aid interpretability

Applied

Fitting Classification Trees

```
library(tree)
library(ISLR)
attach(Carseats)
High <- ifelse(Sales <= 8, "No", "Yes") #recode into binary variable
Carseats <- data.frame(Carseats, High)


# fit classification tree to predict High using all variables but sales
tree.carseats <- tree(High ~ . - Sales, Carseats)
summary(tree.carseats)

##
## Classification tree:
## tree(formula = High ~ . - Sales, data = Carseats)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Income" "CompPrice" "Population"
## [6] "Advertising" "Age" "US"
## Number of terminal nodes: 27
## Residual mean deviance: 0.4575 = 170.7 / 373
## Misclassification error rate: 0.09 = 36 / 400
```

summary function lists variables that are used as internal nodes in tree, the number of terminal nodes, and training error rate

Now we plot the tree:

```
plot(tree.carseats)
text(tree.carseats, pretty = 0)
```



Typing in the tree object prints out the split criterions, # of observations in the branch, deviance, the overall prediction for that branch, and the fraction of observations in that branch that take on values of Yes and No respectively.

```
tree.carseats
```

```
## node), split, n, deviance, yval, (yprob)
##      * denotes terminal node
##
##  1) root 400 541.500 No ( 0.59000 0.41000 )
##    2) ShelveLoc: Bad,Medium 315 390.600 No ( 0.68889 0.31111 )
##      4) Price < 92.5 46 56.530 Yes ( 0.30435 0.69565 )
##        8) Income < 57 10 12.220 No ( 0.70000 0.30000 )
##          16) CompPrice < 110.5 5 0.000 No ( 1.00000 0.00000 ) *
##          17) CompPrice > 110.5 5 6.730 Yes ( 0.40000 0.60000 ) *
##          9) Income > 57 36 35.470 Yes ( 0.19444 0.80556 )
##            18) Population < 207.5 16 21.170 Yes ( 0.37500 0.62500 ) *
##            19) Population > 207.5 20 7.941 Yes ( 0.05000 0.95000 ) *
##        5) Price > 92.5 269 299.800 No ( 0.75465 0.24535 )
##          10) Advertising < 13.5 224 213.200 No ( 0.81696 0.18304 )
##            20) CompPrice < 124.5 96 44.890 No ( 0.93750 0.06250 )
##              40) Price < 106.5 38 33.150 No ( 0.84211 0.15789 )
##                80) Population < 177 12 16.300 No ( 0.58333 0.41667 )
##                  160) Income < 60.5 6 0.000 No ( 1.00000 0.00000 ) *
##                  161) Income > 60.5 6 5.407 Yes ( 0.16667 0.83333 ) *
##                81) Population > 177 26 8.477 No ( 0.96154 0.03846 ) *
##              41) Price > 106.5 58 0.000 No ( 1.00000 0.00000 ) *
```

```

##      21) CompPrice > 124.5 128 150.200 No ( 0.72656 0.27344 )
##      42) Price < 122.5 51 70.680 Yes ( 0.49020 0.50980 )
##      84) Shelveloc: Bad 11 6.702 No ( 0.90909 0.09091 ) *
##      85) Shelveloc: Medium 40 52.930 Yes ( 0.37500 0.62500 )
##      170) Price < 109.5 16 7.481 Yes ( 0.06250 0.93750 ) *
##      171) Price > 109.5 24 32.600 No ( 0.58333 0.41667 )
##      342) Age < 49.5 13 16.050 Yes ( 0.30769 0.69231 ) *
##      343) Age > 49.5 11 6.702 No ( 0.90909 0.09091 ) *
##      43) Price > 122.5 77 55.540 No ( 0.88312 0.11688 )
##      86) CompPrice < 147.5 58 17.400 No ( 0.96552 0.03448 ) *
##      87) CompPrice > 147.5 19 25.010 No ( 0.63158 0.36842 )
##      174) Price < 147 12 16.300 Yes ( 0.41667 0.58333 )
##      348) CompPrice < 152.5 7 5.742 Yes ( 0.14286 0.85714 ) *
##      349) CompPrice > 152.5 5 5.004 No ( 0.80000 0.20000 ) *
##      175) Price > 147 7 0.000 No ( 1.00000 0.00000 ) *
##  11) Advertising > 13.5 45 61.830 Yes ( 0.44444 0.55556 )
##  22) Age < 54.5 25 25.020 Yes ( 0.20000 0.80000 )
##  44) CompPrice < 130.5 14 18.250 Yes ( 0.35714 0.64286 )
##  88) Income < 100 9 12.370 No ( 0.55556 0.44444 ) *
##  89) Income > 100 5 0.000 Yes ( 0.00000 1.00000 ) *
##  45) CompPrice > 130.5 11 0.000 Yes ( 0.00000 1.00000 ) *
##  23) Age > 54.5 20 22.490 No ( 0.75000 0.25000 )
##  46) CompPrice < 122.5 10 0.000 No ( 1.00000 0.00000 ) *
##  47) CompPrice > 122.5 10 13.860 No ( 0.50000 0.50000 )
##  94) Price < 125 5 0.000 Yes ( 0.00000 1.00000 ) *
##  95) Price > 125 5 0.000 No ( 1.00000 0.00000 ) *
##  3) Shelveloc: Good 85 90.330 Yes ( 0.22353 0.77647 )
##  6) Price < 135 68 49.260 Yes ( 0.11765 0.88235 )
##  12) US: No 17 22.070 Yes ( 0.35294 0.64706 )
##  24) Price < 109 8 0.000 Yes ( 0.00000 1.00000 ) *
##  25) Price > 109 9 11.460 No ( 0.66667 0.33333 ) *
##  13) US: Yes 51 16.880 Yes ( 0.03922 0.96078 ) *
##  7) Price > 135 17 22.070 No ( 0.64706 0.35294 )
##  14) Income < 46 6 0.000 No ( 1.00000 0.00000 ) *
##  15) Income > 46 11 15.160 Yes ( 0.45455 0.54545 ) *

```

Now we split train/test to evaluate test error:

```

set.seed(1)
train <- sample(1:nrow(Carseats),200) # 50/50 split
Carseats.test <- Carseats[-train,]
High.test <- High[-train]
tree.carseats <- tree(High~., -Sales, Carseats,subset = train)
tree.pred <- predict(tree.carseats,Carseats.test,type = "class")
table(tree.pred,High.test)

```

```

##      High.test
## tree.pred No Yes
##      No 84 37
##      Yes 35 44

```

correct predictions:


```
(84+44)/200
```

```
## [1] 0.64
```

So the prediction accuracy is around 64%

Next we consider if pruning helps our accuracy:

```
set.seed(1)
cv.carseats <- cv.tree(tree.carseats,FUN = prune.misclass) # pruning our tree
names(cv.carseats)
```



```
## [1] "size" "dev" "k" "method"
```

```
cv.carseats
```

```
## $size
## [1] 20 18 10 8 6 4 2 1
##
## $dev
## [1] 66 66 59 56 53 58 75 85
##
## $k
## [1] -Inf 0.0 0.5 1.5 2.0 4.0 12.0 19.0
##
## $method
## [1] "misclass"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```

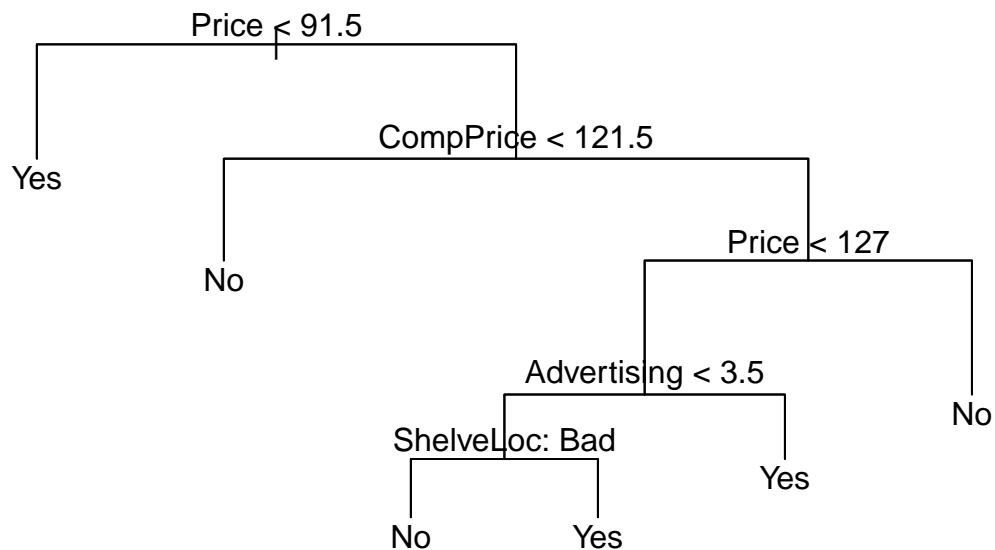
dev is the CV error rate, and thus it looks like the tree with 6 terminal nodes results in lowest error rate, with 53 CV-errors. k corresponds to the cost-complexity parameter used, in which this case was α .

```
par(mfrow=c(1,2))
plot(cv.carseats$size,cv.carseats$dev, type = "b")
plot(cv.carseats$k,cv.carseats$dev,type = "b")
```



Now apply `prune.misclass()` to prune tree to obtain the 6 node tree:

```
prune.carseats <- prune.misclass(tree.carseats, best=6)
plot(prune.carseats)
text(prune.carseats, pretty = 0)
```



Now lets test it on the test data:

```
tree.pred <- predict(prune.carseats,Carseats.test, type = "class")
table(tree.pred,High.test)
```

```
##           High.test
## tree.pred No  Yes
##           No  86  32
##           Yes 33  49
```

```
(86+49)/200
```

```
## [1] 0.675
```

Seems like we have improved on the original tree by about 4%. Not only is it more accurate, but much more interpretable as well.

Fitting Regression Trees


Let's create a train/test split on the Boston data set

```
library(MASS)
set.seed(1)
train <- sample(1:nrow(Boston), nrow(Boston)/2) # 50/50 split
tree.boston <- tree(medv~., Boston, subset = train)
summary(tree.boston)
```

```
##
## Regression tree:
## tree(formula = medv ~ ., data = Boston, subset = train)
## Variables actually used in tree construction:
## [1] "rm" "lstat" "crim" "age"
## Number of terminal nodes: 7
## Residual mean deviance: 10.38 = 2555 / 246
## Distribution of residuals:
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
## -10.1800 -1.7770 -0.1775  0.0000  1.9230 16.5800
```

In regression trees, deviance is the sum of squared errors for the tree

```
plot(tree.boston)
text(tree.boston,pretty=0)
```



It seems that the most important variables are rm(avg number of rooms per dwelling), and lstat(% of lower status of population).

Let's prune the tree:

```
set.seed(1)
cv.boston <- cv.tree(tree.boston)
cv.boston
```

```
## $size
```

```
## [1] 7 6 5 4 3 2 1
##
## $dev
## [1] 4336.868 4321.549 5070.107 5852.631 6560.984 9802.545 19697.191
##
## $k
## [1] -Inf 203.9641 637.2707 796.1207 1106.4931 3424.7810 10724.5951
##
## $method
## [1] "deviance"
##
## attr(,"class")
## [1] "prune" "tree.sequence"
```


```
plot(cv.boston$size,cv.boston$dev, type="b")
```



Using CV, seems like the one with the best performance is the one with 6 terminal nodes


Now let's prune the tree:

```
prune.boston <- prune.tree(tree.boston,best = 6)
plot(prune.boston)
text(prune.boston,pretty=0)
```



Let's now make predictions on test set:

```
yhat <- predict(prune.boston,newdata = Boston[-train,])
boston.test <- Boston[-train,"medv"]
plot(yhat, boston.test)
abline(0,1)
```



```
mean((yhat - boston.test)^2)
```

```
## [1] 35.16439
```

So the test set MSE is around 35.

Baggin and Random Forests

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
set.seed(1)
```

```
# Bagging
```


```
bag.boston <- randomForest(medv~.,
                           data = Boston,
                           subset=train,
                           mtry = 13,
                           importance = T) # we set mtry = 13 since we are using all predictors
```

```
bag.boston
```

```
##
## Call:
## randomForest(formula = medv ~ ., data = Boston, mtry = 13, importance = T, subset = train)
##           Type of random forest: regression
##           Number of trees: 500
## No. of variables tried at each split: 13
##
##           Mean of squared residuals: 11.39601
##           % Var explained: 85.17
```

Let's see how well bagging performs on test set:

```
yhat.bag <- predict(bag.boston,newdata = Boston[-train,])
plot(yhat.bag,boston.test)
abline(0,1)
```



```
mean((yhat.bag-boston.test)^2)
```

```
## [1] 23.59273
```

MSE is around 23.6, which is significantly better than the optimally-pruned tree from before

Let's try using random forest with $m = 6$:


```

set.seed(1)
rf.boston <- randomForest(medv~.,
                          data = Boston,
                          subset = train,
                          mtry = 6,
                          importance = T)
yhat.rf <- predict(rf.boston, newdata = Boston[-train,])
mean((yhat.rf-boston.test)^2)

```

```
## [1] 19.62021
```

we can see that the test MSE has decreased down to 19.6, which is an improvement over bagging

We can see the importance of each variable using the `importance()` function:


```
importance(rf.boston)
```

##		%IncMSE	IncNodePurity
##	crim	16.697017	1076.08786
##	zn	3.625784	88.35342
##	indus	4.968621	609.53356
##	chas	1.061432	52.21793
##	nox	13.518179	709.87339
##	rm	32.343305	7857.65451
##	age	13.272498	612.21424
##	dis	9.032477	714.94674
##	rad	2.878434	95.80598
##	tax	9.118801	364.92479
##	ptratio	8.467062	823.93341
##	black	7.579482	275.62272
##	lstat	27.129817	6027.63740

The left column %IncMSE is based on the mean decrease of accuracy in predictions on the OOB samples when given variable is excluded from the model. IncNodePurity is a measure of total decrease in node purity that results from splits over that variable, averaged over all trees. We can plot these below:

```
varImpPlot(rf.boston)
```

rf.boston



We can see that rm and lstat are by far the two most important variables

Boosting


```
library(gbm)
```

```
## Loaded gbm 2.1.8
```

```
set.seed(1)
```

```
boost.boston <- gbm(medv~.,
  data = Boston[train,],
  distribution = "gaussian",
  n.trees = 5000,
  interaction.depth = 4) # would use dist = bernoulli if it were classification
```

```
summary(boost.boston)
```




```
##      var    rel.inf
## rm      rm 43.9919329
## lstat   lstat 33.1216941
## crim    crim 4.2604167
## dis     dis 4.0111090
## nox     nox 3.4353017
## black   black 2.8267554
## age     age 2.6113938
## ptratio ptratio 2.5403035
## tax     tax 1.4565654
## indus   indus 0.8008740
## rad     rad 0.6546400
## zn      zn 0.1446149
## chas    chas 0.1443986
```


Again, we can see that `rm` and `lstat` are by far the most important variables

We can produce a *partial dependence plot* for these two variables, which illustrate the marginal effect of the variables on the response after *integrating out* the other variables:

```
par(mfrow = c(2,2))
plot(boost.boston,i="rm")
```



```
plot(boost.boston,i="lstat")
```



We can see that house prices are increasing with rm and decreasing with lstat

Finally, we can use the boosted model to fit the test data:

```
boost.boston <- gbm(medv~.,
  data = Boston[train,],
  distribution = "gaussian",
  n.trees = 5000,
  interaction.depth = 4,
  verbose = F,
)

yhat.boost <- predict(boost.boston,
  newdata = Boston[-train,],
  n.trees = 5000)
mean((yhat.boost - boston.test)^2)
```

```
## [1] 18.18255
```

The shrinkage default for `gbm()` is 0.1. We can see an incremental improvement of using boosting over random forests (improvement by about 1%).